When looking at a section of host data 75 days in the past, what interval of granularity can you see?

35 days

5 mins

1 min

<mark>1 hr</mark>

By default, how often does dynatrace auto-discover, analyse and store log files

Every 2 hrs

Every 15 min

Every 24hrs

Every 60 sec

Within dynatrace where do tags and metadata come from?

Tags are detected during startup or discovery of a monitored entity. Metadata is created manually or via rules

Both tags and metadata can be added manually and created via rules in dynatrace

<mark>Metadata is detected during startup or discovery of a monitored entity. Tags are created manually, imported from cloud or via rules</mark>

Tags and metadata are added manually to entities as needed

How does dynatrace perform evaluations for rapidly increasing response time degradation

Evaluation are performed based on sliding 5 min time intervals

Evaluations are based on the slowest 10% of all requests during the degradation

Evaluations are based on the number of users impacted by the response time degradation. The sliding timeframe for evaluation is increased by 5 mins for every 500 users affected

Evaluations are performed based on sliding 15 min time intervals

Which of the following are the main functions of an active gate? Select all that apply

None of the above

Act as entry points for sealed networks, Buffering and compression of oneagent messages and delivery to the Dynatrace cluster

Detection of external services, such as database services for collection of performance

Authentication of Oneagent requests

When Synthetic monitor is executed, which browser is used? Select all that apply

Correct answer: Chrome, Microsoft Edg

Why might there be an Opaque service in your environment? Select all that apply

The service technology is supported by dynatrace but code-level visibility isn't available, The service technology is unrecognised or unsupported

Code-level visibility isn't possible because the services have external dependencies

Services have been aggregated because they are of the same process group and perform identical functions across separate cluster nodes

All of the above

How can you track which users in an application have successfully completed a checkout?

By Uploading a source map file

By editing your content capture settings

By Setting a conversion goal

By Setting a user tagging rule

How can you show the locations of users with Private IP addresses on the world map

You can manually place the user locations on the World Map

By re-ordering the detection rules on the IP detection page

by Configuring the Geolocation breakdown table

By Mapping IP addresses to locations in the web & Mobile monitoring settings

Where can you view the attributes for service, such as the web server name and web application ID?

In the PROPERTIES and TAGS section of the Service overview page

In the Transactions & Services View

In the attributes section of the service overview page

Via the service View Metadata action

Queston: 1Which of he following s a emen s abou user session propertes is rue?A.You can use single users session proper y access multple applicatonsB.User session propertes can be used in USQL queriesC.User session propertes are available for web and mobile applicatonsD.User session propertes are re rospectve. Answer: B, CQueston :2When analyzing a service ow or dis ribu ed race, wha is he processing tme calcula ed as?A.The tme be ween he s ar tme and he end tme of he ransactonB.The tme be ween when he clien riggers he reques and receives he responseC.The o al tme is akes he CPU o execu e code.D.The sum of he tme aken by all asynchronous executons occurring during he ransactonAnswer: A

Queston: 3When you need o use he RUM browser ex ension?A.Any tme you wan o rack real user moni oring da aB.I is required o record a browser clickpa hC.When you wan real user moni oring bu don' have access o he underlying applicaton server or he HTML page sourceD.When you wan real user moni oring and have access he applicaton server bu don' have access o he HTMLpage source.Answer: CQueston: 4When is a Ne work Zone in Dyna race?A.A de ned group ha provides an ex ra layer of con ex ual informaton abou your ne work moni oring da aB.A de ned group ha allows you o perform bulk operatons across multple ne works moni ored by Dyna raceC.An ent y ha represen s your ne work s ruc ure and is used o e cien ly rou e ra c for OneAgen and actve ga es.D.An informaton-parttoning

mechanism used for l er da a for speci c eams.Answer: CQueston: 5How long does Dyna race keep tmeseries me rics?A.35 daysB.10 daysC.5 yearsD.30 daysAnswer: cQueston: 6

What are the key capabilities of the Dynatrace platform?

Observability, security, automation, integration, and business analytics.

What does the Dynatrace platform provide for observability?

Real-time monitoring and visualization of application performance and infrastructure.

What is the role of AI in the Dynatrace platform?

AI automates discovery, analysis, and key business processes, providing precise answers and intelligent automation.

What is DQL in the context of Dynatrace?

DQL stands for Dynatrace Query Language, which is used to analyze data within the platform.

What is the Dynatrace Hub?

A marketplace for apps, extensions, and technologies that enhance the platform's functionality.

What kind of support does the Dynatrace Support Center provide?

Documentation, community discussions, on-demand courses, live training, and certifications.

What is the goal of the Dynatrace platform regarding business outcomes?

To predict, prevent, and resolve issues before they impact customers.

What is the estimated time to navigate the Dynatrace User Interface?

3 minutes.

What is the estimated time to read the Dynatrace documentation on navigating the platform?

10 minutes.

What types of technologies does Dynatrace integrate with?

Cloud, container, and various technology integrations.

What is the primary focus of Dynatrace's platform?

To maintain system resilience and ensure optimal digital experiences.

How does Dynatrace collect data for monitoring?

By collecting and analyzing data from every layer of the technology stack, from user interactions to backend processes.

What is the benefit of AI-powered insights in Dynatrace Apps?

They help reduce downtime and improve performance, ensuring high-quality user experiences.

What are Dynatrace Notebooks used for?

For custom exploration and visualization of observability data, integrating external sources and predictive analytics.

What is the role of Dashboards in Dynatrace?

To offer visual interfaces that provide real-time insights into application and infrastructure performance.

How do Notebooks and Dashboards differ in Dynatrace?

Notebooks allow for custom exploration of data, while Dashboards aggregate key metrics for monitoring and analysis.

What is Business Analytics in Dynatrace?

A feature that connects application performance and user experience to business metrics for collaborative insights.

What are Business Events in the context of Dynatrace?

Key components that link application performance and user experience to business metrics, aiding in decision-making.

What is Digital Experience Monitoring (DEM)?

A feature that provides a unified interface for tracking and analyzing key performance indicators and user experience metrics.

What is the significance of Real User Monitoring (RUM) in Dynatrace?

It helps track real user interactions with applications to assess performance and user experience.

What is Synthetic Monitoring in Dynatrace?

A method to simulate user interactions to monitor application performance and availability.

What is the role of predictive analytics in Dynatrace Notebooks?

To enhance data exploration and visualization by anticipating trends and issues.

How do dashboards facilitate decision-making in Dynatrace?

By providing a centralized hub for monitoring, analyzing, and troubleshooting key performance metrics.

What is the significance of integrating external sources in Dynatrace Notebooks?

It allows for a comprehensive view of observability data and enhances analysis capabilities.

What do Business Events help monitor?

User interactions, track performance, and identify trends.

What are the key components of Digital Experience Monitoring?

Real User Monitoring (RUM) and Synthetic Monitoring.

What does Real User Monitoring (RUM) track?

Real-time user interactions within applications.

How does RUM contribute to Business Analytics?

It provides insights into user behavior and performance metrics.

What is the purpose of Synthetic Monitoring?

To simulate user interactions with applications using automated scripts to identify performance issues.

What insights do User Analytics provide?

Data on user interactions, including metrics like bounce rates, top landing pages, exit pages, and conversion rates.

What is session segmentation in user analytics?

Tracking and analyzing a user's journey by breaking interactions into smaller parts.

What does the Performance Analysis section display?

Various performance metrics for applications, helping to monitor key metrics and identify issues.

What is the purpose of Dashboard Reports?

To visualize and analyze data from applications, offering interactive visualizations to explore trends and identify root causes.

What does the Dashboards API facilitate?

Seamless data analysis across different platforms by integrating with other tools.

What is DQL (Dynatrace Query Language)?

A tool for querying, analyzing, and visualizing performance data stored in Dynatrace Grail storage.

What advantages does DQL offer over traditional databases?

It requires no up-front description of the input data's schema, unlike SQL tables.

What can users do with DQL?

Spot trends, find anomalies, perform statistical modeling, and understand system behavior.

What is the role of Davis AI in the Dynatrace Platform?

To link events and provide context, making it easier to find the root cause of problems.

What types of metrics can be tracked in the Dynatrace Platform?

Response times, error rates, and resource use.

How can users explore RUM data?

By accessing it in various ways, including Notebooks and Dashboards.

What does Dynatrace highlight when clicking into a problem in the Problems app?

It provides automatic detection of issues.

What is the purpose of the 'fetch bizevents' query?

To return data related to business events.

What is the benefit of creating custom dashboards in Dynatrace?

To explore trends and compare performance over time.

How can users subscribe to Dashboard Reports?

By visiting Dynatrace Docs and following the subscription process.

What does the Davis Security Advisor do?

It recommends fixes to improve the overall security of the environment.

What are third-party vulnerabilities?

Weaknesses in external vendors' systems that can compromise an organization's security.

What are common sources of third-party vulnerabilities?

Software flaws, inadequate security protocols, and lack of visibility into vendor practices.

What are code-level vulnerabilities?

Flaws in software code that can be exploited to compromise security.

What is the purpose of the Security Investigator in Dynatrace?

To streamline evidence-driven investigations using logs, metrics, and traces.

What does the Vulnerabilities feature in Dynatrace help with?

Detecting, visualizing, analyzing, monitoring, and fixing vulnerabilities in the application stack.

What is the significance of ingestion in the Dynatrace Platform?

It facilitates the acquisition and integration of critical data streams from diverse sources.

What is log ingestion?

The process of collecting log data from various sources within an infrastructure.

What does the VMware Monitor Hub do?

Enables the ingestion of metrics and logs from VMware vSphere environments.

What does Dynatrace Topology visualize?

It automatically discovers and maps all components and dependencies in the application environment.

How does Dynatrace assist in log management and analytics?

It provides direct access to log content and allows for narrowing down log records for analysis.

What is the purpose of the Dynatrace Hub?

To enhance software development and deployment through automated observability and IT operations management.

What is the function of the Logs Ingestion API?

To collect log data in environments where OneAgent installation isn't possible.

What is the function of the pipe character '|' in DQL?

The pipe character connects commands, sending the output from one command to the input of the next, enabling powerful data processing.

What command is used to load data in DQL?

The fetch command is used to specify the type of data to load, such as logs, entities, events, or metrics tables.

What does the fetch command syntax look like?

An example syntax is: fetch logs, from: now()-25m, to: now() | limit 10.

What is the purpose of the filter clause in DQL?

The filter clause refines results based on boolean conditions, allowing users to focus on specific timeframes or application tiers.

How can you filter logs for a specific log level in DQL?

You can use: fetch logs | filter loglevel == 'INFO'.

What is the purpose of the fields command in DQL?

The fields command specifies which data to retrieve and allows renaming existing fields or adding new fields.

How do you sort results in DQL?

You can sort results using the sort command, for example: fetch logs | sort loglevel asc.

What is the summarize command used for in DQL?

The summarize command performs calculations on a set of values and returns a single value, similar to SQL's GROUP BY clause.

What is an example of using the summarize command?

An example is: fetch logs | summarize count(), by: {status}.

What command is used to parse and extract information from data in DQL?

The parse command is used to extract components of a string content into distinct fields.

What is an example of using the parse command?

An example is: record(content='SourceTable=AAA, RowCount=1001, TargetTable=XYZ') | parse content, "('SourceTable=' LD:source ', RowCount=' INT:rows ', TargetTable=' LD:target)".

What are the three filtering commands available in DQL?

The three filtering commands are filter, dedup, and filterOut.

How can you exclude records with a specific log level in DQL?

You can use: fetch logs | filterOut log.level == 'INFO'.

What is the role of the limit command in DQL?

The limit command restricts the number of records returned in the results.

What does the in function do in DQL?

The in function filters logs based on whether a field matches any value in a specified array.

What is the equivalent of SQL's SELECT statement in DQL?

In DQL, the equivalent is the fetch command, which retrieves data from specified sources.

What is the significance of the dt.entity.* syntax in DQL?

This syntax is used to query from monitored entities presented as views in Dynatrace.

How can you rename fields in DQL?

You can rename fields using the fields command, for example: fields NewFieldName = oldFieldName.

What is the benefit of using the parse command on log records?

It allows for structured queries by extracting specific components into distinct fields.

What is a full example of a DQL query that summarizes data?

An example is: fetch logs | filter log.source == 'My Log Source' and contains(content,'HTTP_POST') | fields timestamp, content | parse content, 'LD 'tenantId=' WORD:tenantId ', duration=' INT:duration' | sort duration desc | summarize sum(duration), by: tenantId.

What is the first step in preparing for Dynatrace installation?

Determine what services or processes you want to monitor.

What are the different monitoring environments you might configure?

Production clusters, development environments, and staging servers.

How do you start the installation of OneAgent?

Select the 'Get data into Dynatrace' icon from your first host, then download and install OneAgent.

What types of entities can OneAgent monitor?

Servers, applications, services, databases, and more.

What is RUM in Dynatrace?

Real User Monitoring is available through agentless monitoring, PaaS integration, and browser extensions.

What does synthetic monitoring do?

Simulates user visits and monitors the availability of your application.

What is the purpose of ActiveGates in Dynatrace?

To facilitate communication between OneAgent and the Dynatrace server, especially in secure environments.

What operating systems does Dynatrace support for installation?

z/OS, Windows, Solaris, Linux, and AIX.

Which cloud platforms can Dynatrace run on?

AWS, Heroku, Azure, and Google Cloud.

What are the three modes of OneAgent?

Full-stack monitoring mode, full-stack + DEM mode, and custom monitoring modes.

What is the primary function of OneAgent?

To collect data within your monitored environment and provide operational and business metrics.

How does OneAgent communicate with Dynatrace?

Communication is outbound only; Dynatrace never initiates communication with OneAgent.

What is the process for deploying OneAgent?

Select the 'set up' icon, choose your platform, and follow the installation instructions.

What is required for OneAgent installation on Windows?

You will be prompted to create a new PaaS token or enter an existing token.

What is an ActiveGate?

A secure proxy between OneAgent and the Dynatrace server that ensures data is transmitted securely.

What happens if you do not use ActiveGate?

Dynatrace functions entirely in the cloud, which may raise security concerns.

What is the role of ActiveGate in monitoring?

To route OneAgent traffic to Dynatrace and monitor cloud environments and remote technologies.

What should you do if you want to monitor a cloud-based installation securely?

Download ActiveGate to your data center.

What is required for container platform installations of OneAgent?

A Dynatrace operator token is needed for Kubernetes or OpenShift installations.

What should you do if you want to use ActiveGate in your environment?

Install ActiveGate before deploying OneAgent to ensure immediate connection.

What does OneAgent automatically do once deployed?

Detects and monitors a wide spectrum of host processes, applications, networks, and health metrics.

How does OneAgent provide Real User Monitoring (RUM)?

By automatically injecting a JavaScript tag into the HTML of your application pages.

What is the significance of the Dynatrace ActiveGate connection?

It allows for secure data collection and transmission to the Dynatrace Cluster.

What is the purpose of installing ActiveGate?

To route z/OS traffic to Dynatrace for mainframe monitoring.

What are the two types of ActiveGate deployments?

Host-based ActiveGate and containerized deployment.

What is the function of a host-based ActiveGate?

It is deployed on a physical or virtual host and can have multiple modules installed.

What is the limitation of a containerized ActiveGate deployment?

It is configured to run only one module.

What does enabling mobile monitoring provide?

Insight into the performance level experienced by mobile users.

What is the first step to deploy Dynatrace for mobile applications?

Select 'open' then 'deploy Dynatrace' and enter the name of the mobile application.

What is Agentless RUM?

A feature that automatically enables monitoring for all detected applications without requiring OneAgent installation.

When should Agentless RUM be used?

When OneAgent cannot be installed.

What is required to monitor a web frontend process?

Check settings under monitoring overview and process groups.

What is the role of personal access tokens?

Configured by and for individual users for secure access.

What is the recommended deployment option for Dynatrace in Kubernetes?

Deploy Dynatrace Operator in cloud-native full stack mode.

What does the Dynatrace operator manage?

Automated rollout, configuration, and lifecycle management of Dynatrace components.

What is the function of the CSI driver in Dynatrace deployment?

Provides OneAgent code modules for application pods while minimizing storage usage.

What is a key benefit of using the CSI driver?

OneAgent requires only read access, improving performance and control over license usage.

What steps are involved in monitoring Kubernetes with Dynatrace?

Select distribution, observability options, enter cluster and network zone, generate API tokens, download YAML, and install using helm or kubectl.

What prerequisites are needed to install Dynatrace via the web UI?

Helm and Kubectl installed, access to a supported Kubernetes cluster, and install OneAgent permissions.

What is the first command to run after applying the YAML file for Dynatrace installation?

Install the Dynatrace operator.

What does the Dynatrace Webhook do?

Modifies pod definitions to include Dynatrace code modules needed for application observability.

What is required for Kubernetes app monitoring in Dynatrace?

A Dynatrace platform subscription and permissions to access the Kubernetes app.

What is the purpose of generating API tokens in Kubernetes app monitoring?

To authenticate and enable monitoring functionalities.

What is the significance of the Dynatrace operator token?

It allows the operator to install ActiveGate and OneAgents to the cluster.

What should be done after the JS tag is generated for Agentless RUM?

Paste the tag into each application's HTML pages to monitor them.

What is the importance of the script tag placement for monitoring?

The script tag must be placed as the first JS file in the page for proper functionality.

What command is used to connect to the cluster for deployment in Dynatrace?

Use the copy command in the terminal where kubectl and Helm are installed.

What is the purpose of Fluent Bit in Kubernetes?

Fluent Bit is used for streaming logs to Dynatrace.

What prerequisites are needed for using Fluent Bit with Dynatrace?

Set up security context constraints (SCC) for OpenShift, use Helm version 3, allow egress traffic from the Fluent Bit namespace to Dynatrace, and use Dynatrace operator version 1.1.0+.

What are the steps to configure Dynatrace for Kubernetes?

Select AKS for distribution, choose observability options, configure the cluster with a name and network zone, provide or generate an operator token, and download dynakube.YAML.

What types of ActiveGates does Dynatrace support?

Dynatrace supports Environment ActiveGate and Cluster ActiveGate.

What is the difference between Environment ActiveGate and Cluster ActiveGate?

Environment ActiveGate is used for Dynatrace SaaS, while Cluster ActiveGate is installed in a network segment accessible by all agents.

What is the function of Cloud-native ActiveGate?

Cloud-native ActiveGate is packaged in a cloud container and only supports a single module.

What are the three purposes of ActiveGates?

1. Route OneAgent traffic to Dynatrace and monitor cloud environments. 2. Run Synthetic monitoring from a private dedicated location. 3. Route z/OS traffic to Dynatrace for Mainframe monitoring.

What is the Grail Data Model in Dynatrace?

The Grail Data Model consists of data ingestion, processing, storage, and querying using DQL.

What are Grail Buckets used for?

Grail Buckets are used to set retention times, store data, and manage user permissions.

What types of buckets exist in Dynatrace?

There are built-in buckets created by Dynatrace and custom buckets created by users.

How many custom buckets can be created per Dynatrace environment?

Up to 80 custom buckets.

What is the IAM framework used in Dynatrace for access control?

The IAM framework is based on an Attribute-Based Access Control (ABAC) protocol.

What is the purpose of configuring network zones in Dynatrace?

Network zones represent the network structure and should contain ActiveGates accessible by OneAgent and zRemote without a firewall in between.

What is the process for setting up process groups in Dynatrace?

Process groups are set up by detecting processes that perform the same function across one or more hosts and customizing detection rules.

What are clustered services in Dynatrace?

Clustered services refer to the same service running in multiple processes within the same process group.

What is the role of custom services in Dynatrace?

Custom services can have multiple entry points and are recorded by OneAgent for monitoring.

What is a key request in Dynatrace?

Key requests are mission-critical requests that have advanced monitoring capabilities and trend lines retained indefinitely.

How can you mark a request as a key request in Dynatrace?

Select the service, view key requests, and mark the desired request as a key request.